



UNIVERSITY OF  
CENTRAL FLORIDA



# ALISA: Accelerating Large Language Model Inference via Sparsity-aware KV Caching

---

Youpeng Zhao, Di Wu, Jun Wang

*Computer Systems and Data Science (CASS) Lab,  
University of Central Florida*



# Background

## - Emergence of Large Language Models (LLMs)

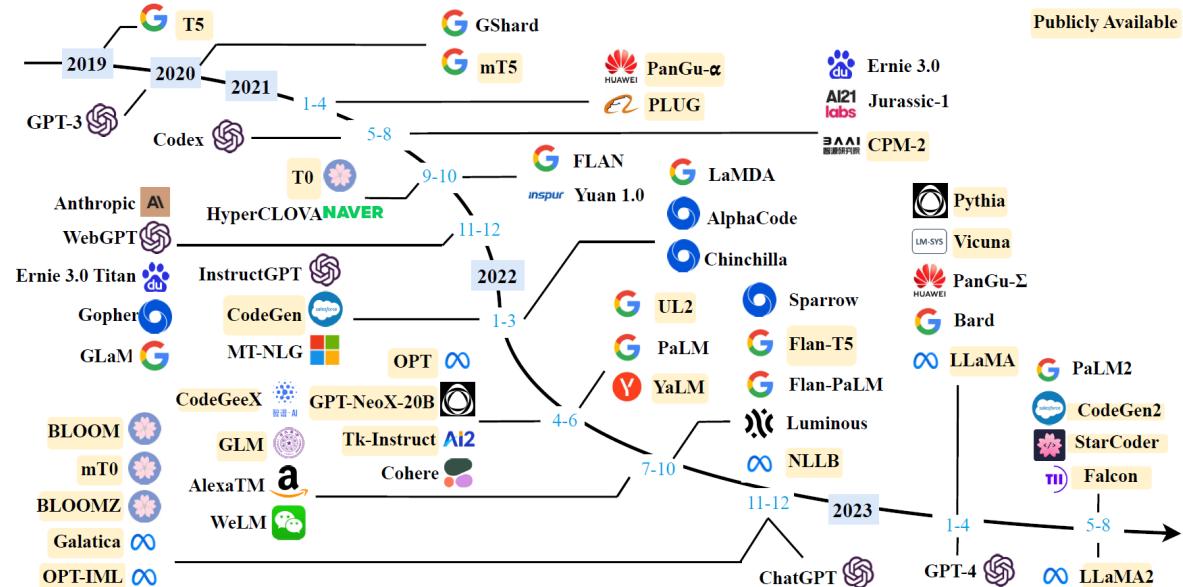
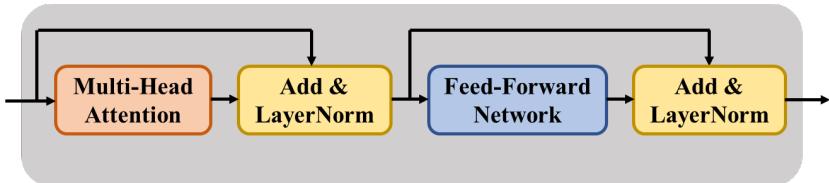


Fig. 1: A timeline of existing large language models (>10 B) in recent years [1].

- Inference serving accounts for most LLM-based application scenarios. Accelerating LLM inference has become an increasingly important research problem.

# Background

## - Transformer Architecture



$$AW(Q, K) = \sigma\left(\frac{QK^T}{\sqrt{d}}\right)$$

$$Attn(Q, K, V) = AW(Q, K) \cdot V$$

**Quadratic Complexity**

## - LLM Inference with KV Caching

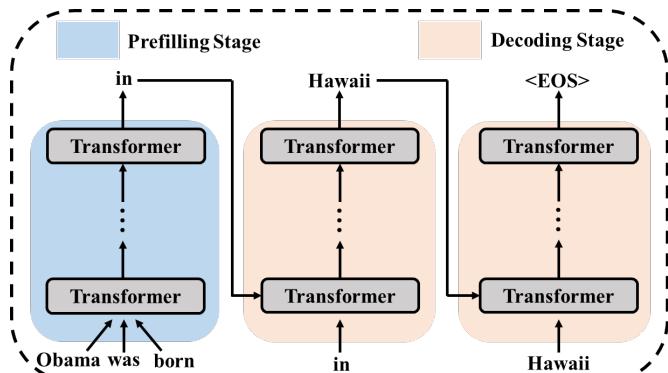


Fig. 2: Autoregressive inference of LLMs

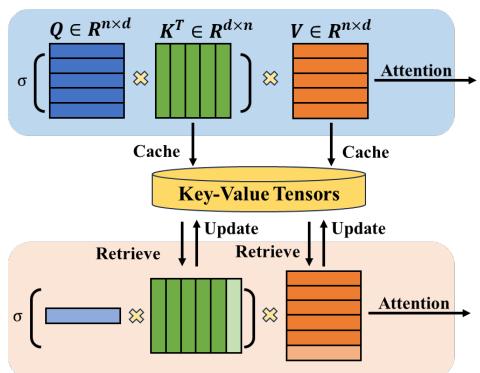
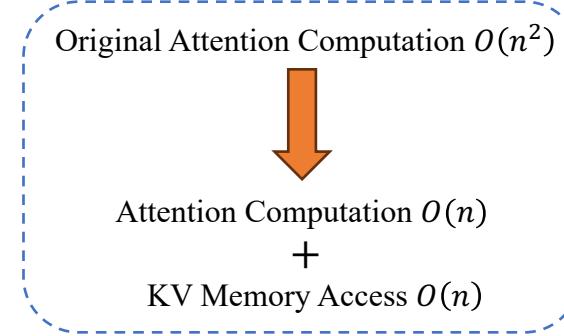


Fig. 3: KV Caching Mechanism [2].



# Motivation

- KV caching increases memory overhead

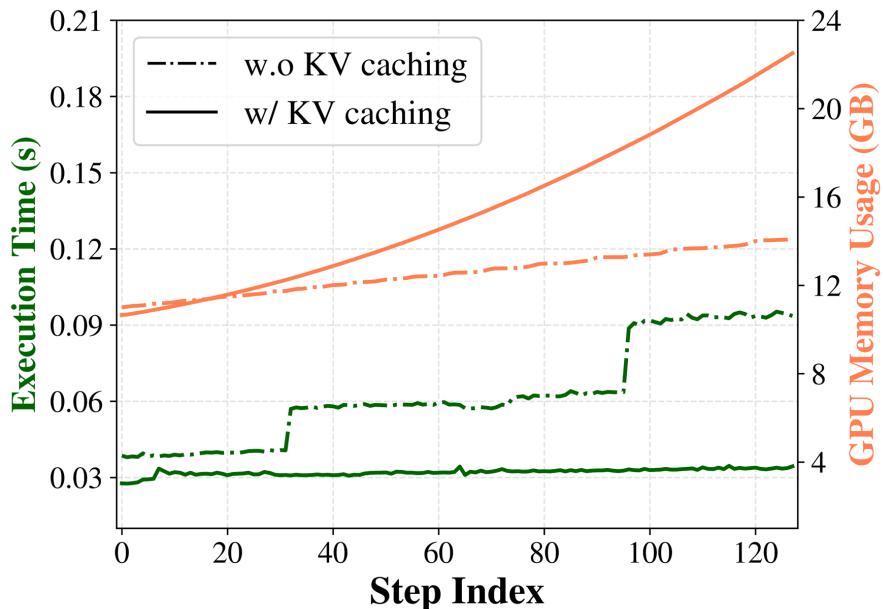
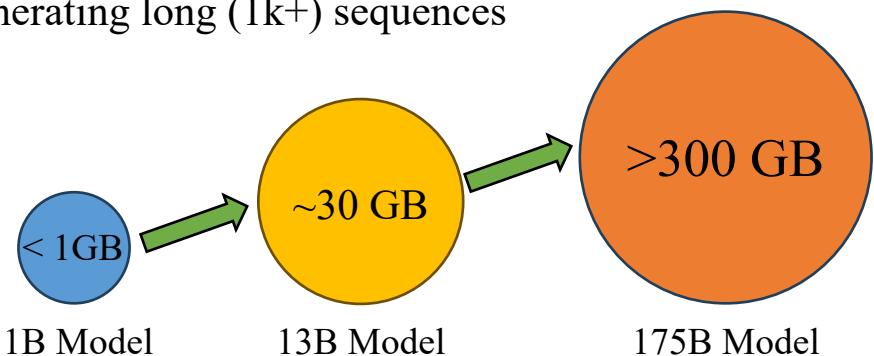


Fig. 4: Execution time and memory comparison for OPT-6.7B.

$$KV\ Tensor\ Size = B \cdot h \cdot n$$

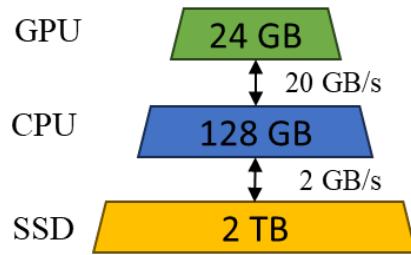
- For small LLMs, the KV tensor size for one token is around dozens of MB; for larger models, the size can be as big as dozens of GB
- The memory footprint of KV tensors can be a potential problem when serving large models or generating long (1k+) sequences



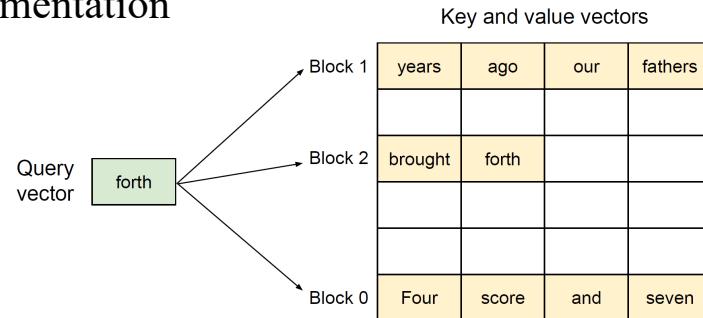
# Motivation

## - Optimizing KV Caching (Previous works)

- FlexGen [3]: Utilize CPU DRAM, and secondary storage (SSD) to hold the intermediate KV tensors and model weights



- vLLM [4]: Employ non-contiguous paged memory to store KV tensors at block-level, where each block contains a fixed group of tokens to reduce memory fragmentation



## - Other Works for accelerating LLM inference:

- Attention computation acceleration [5,6], pruning [7,8] (Algorithm)
- Hardware acceleration [9,10] (Accelerator co-design)
- Quantization [11,12]

**Focus on model weight reduction**

**Cannot scale to large LLMs**

# Main Problem

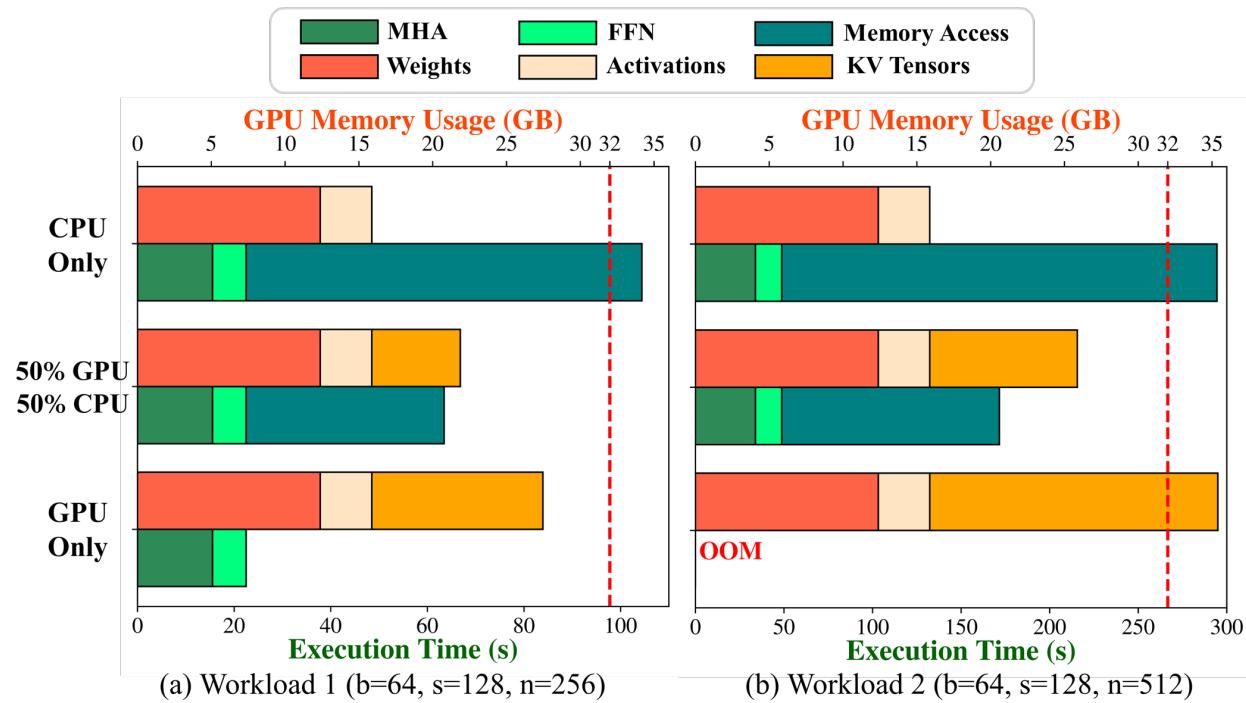


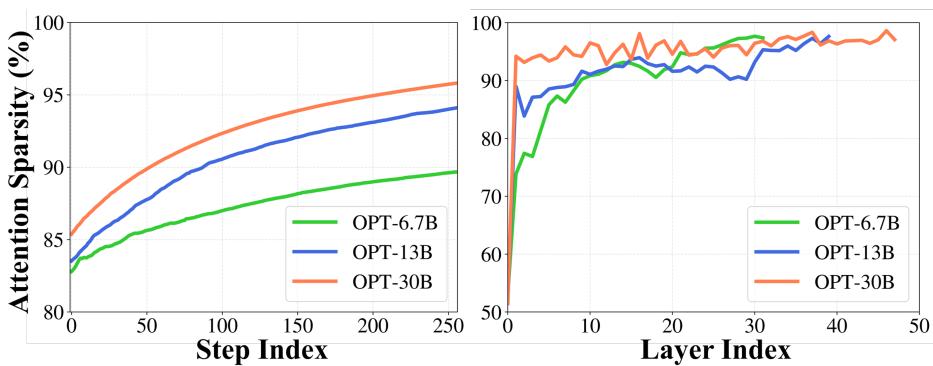
Fig. 5: Execution breakdown of OPT-6.7B in a single NVIDIA V100 (32GB) GPU.

Frequent offloading/reloading incurs significant I/O latency, creating new bottleneck for LLM inference

# Proposal

How do we innovate KV caching to alleviate I/O bottleneck of LLM inference?

**Key Observation: Not all tokens are created equal!**



**Fig. 6: Attention sparsity observed across different steps and layers of OPT model inference.**

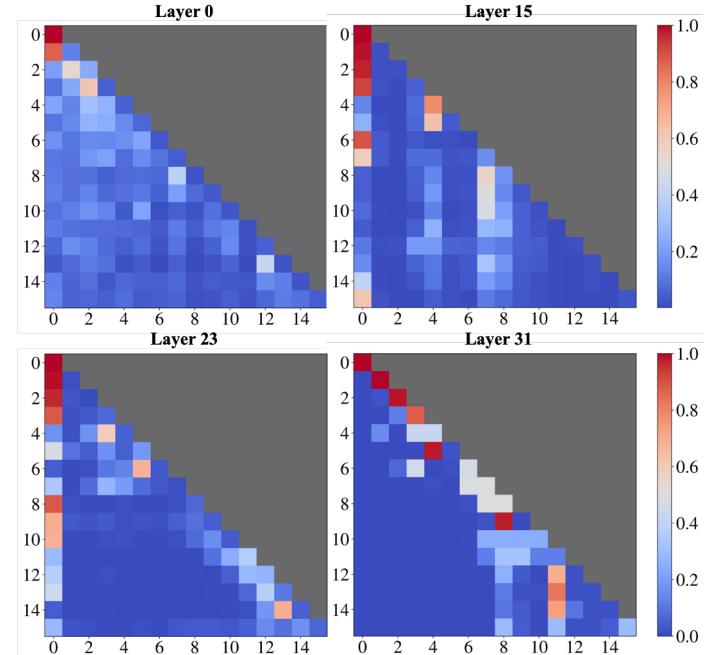
$$A = \begin{bmatrix} A_{11} & 0 & 0 & 0 \\ A_{21} & A_{22} & 0 & 0 \\ A_{31} & A_{32} & A_{33} & 0 \\ A_{41} & A_{42} & A_{43} & A_{44} \end{bmatrix} \quad \text{Attn}_4(Q, K, V) = A_{41}V_1 + A_{42}V_2 + A_{43}V_3 + A_{44}V_4$$

Assume  $A_{42}$  is close to zero, we have

$$\text{Attn}_4(Q, K, V) \approx A_{41}V_1 + A_{43}V_3 + A_{44}V_4$$

Therefore, we do not need to calculate  $A_{42}$ , and do not need KV tensors for the second token

**Fig. 7: Attention weight distribution for OPT-6.7B.**



# Proposal

## How do we leverage sparsity for KV caching?

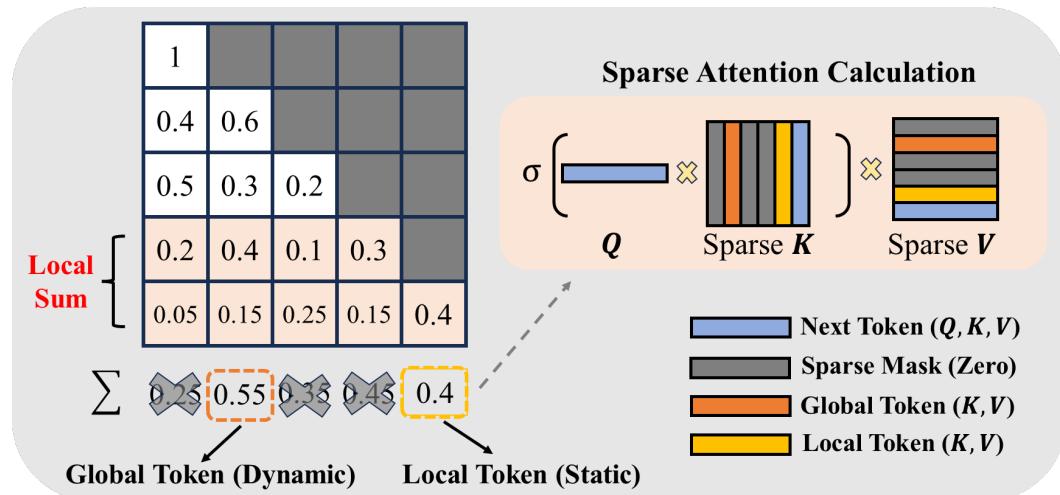
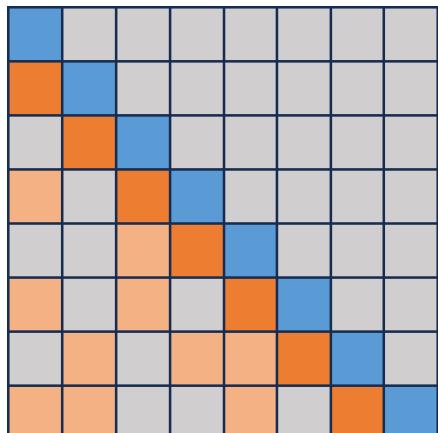
1. Identifying important tokens (Algorithm): we need a low-cost mechanism to distinguish important tokens without hurting LLM accuracy
2. Sparsity-aware Caching Policy (System): when GPU cannot hold all the KV tensors, we need to design a suitable low-overhead caching policy to allocate KV tensors between CPU and GPU and ensure a low miss rate;
3. Caching vs. Recomputation (System): for longer sequences, the benefits of KV caching diminishes, we need to consider recomputation of partial KV tensors instead of caching.

# Algorithm Design

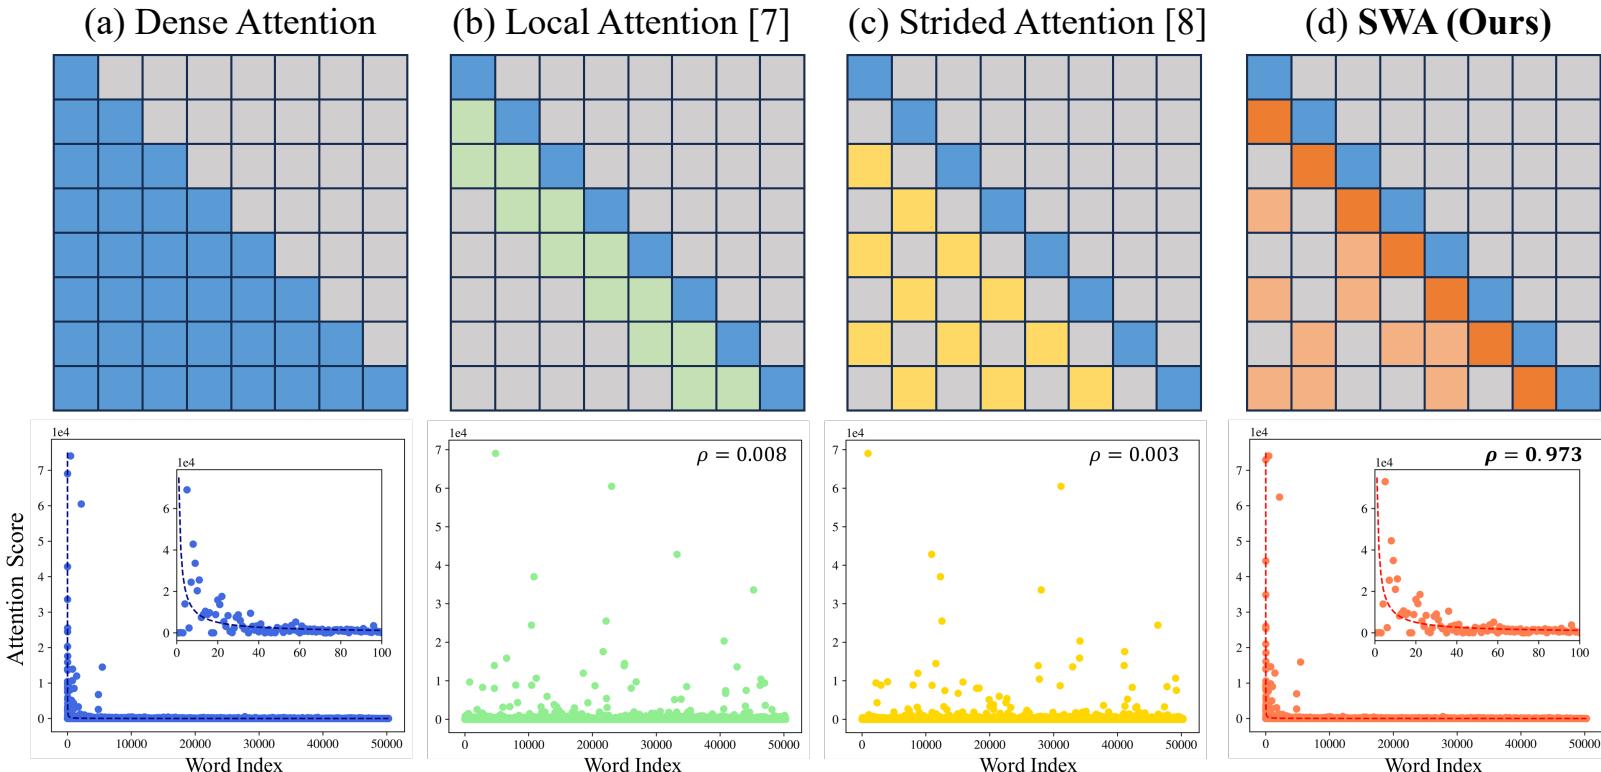
## Sparse Window Attention (SWA):

- Global Dynamic Sparse Patterns: determined by local attention weight sum (light orange color)
- Local Static Sparse Patterns: determined by recency (dark orange color)

**SWA Generated Patterns**



# Algorithm Design



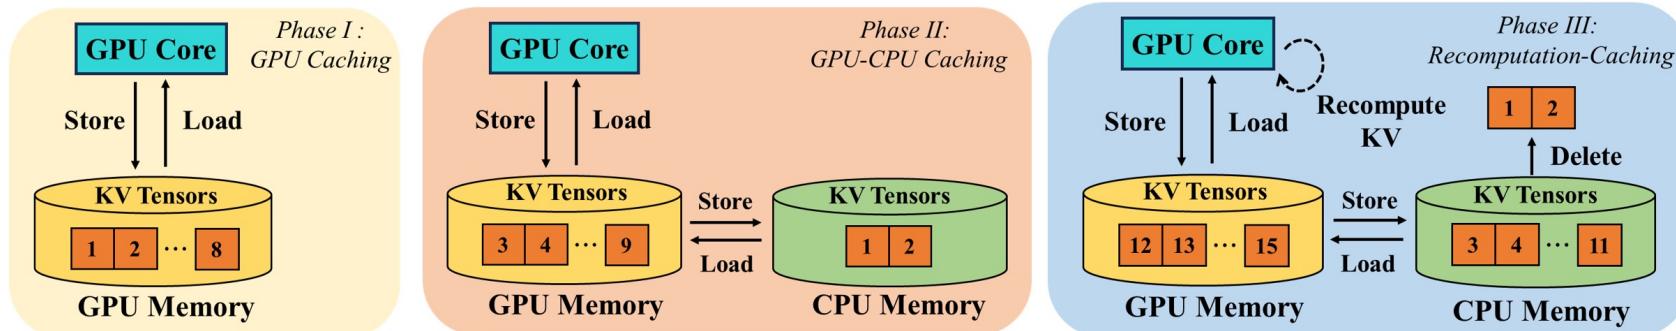
**Fig. 8: Comparison of different attention methods. Top row illustrates the attention patterns, and the bottom row compares the final attention score distribution**

# System Design

## Three-phase Dynamic Scheduling:

We divide the LLM inference into three phases to balance the caching and Recomputation

- Phase I: *GPU Caching*. All KV tensors can fit in GPU memory.
- Phase II: *GPU-CPU Caching*. Split KV tensors at token-level on both GPU and CPU.
- Phase III: *Caching-Recomputation*. Delete partial KV tensors and perform recomputation instead if needed



# System Design

## How do we determine the phase switch step and offload/recomputation ratio?

We formulate this question into an optimization problem to minimize total execution time.

- Size of KV tensors:  $4 \cdot b \cdot l \cdot h$
- Number of tokens moved from GPU to CPU:  $\theta_j^c = \alpha(j + s)$
- Number of tokens moved from GPU to CPU:  $\theta_g^c$

For each step, the execution time is:

$$T_j^m(\alpha) = \frac{4 \cdot b \cdot l \cdot h \cdot (\theta_j^c + \theta_g^c)}{B}$$

The total execution time is:

$$\min_{\{\alpha, \beta, p_1, p_2\}} \sum_{j=1}^{p_2} T_j^c + \sum_{j=p_1}^{p_2} T_j^m(\alpha) + \sum_{j=p_2}^n T_j^r(\beta)$$

TABLE II: Notations.

|                           |  |
|---------------------------|--|
| $h, l, b$                 | hidden dimension, layer count, batch size  |
| $s, n$                    | input length, output length                |
| $r, B$                    | KV caching ratio, CPU-GPU bandwidth        |
| $\alpha, \beta, p_1, p_2$ | offload/recompute ratio, phase switch step |
| $T^c, T^r$                | Time for compute and recompute             |
| $T^m$                     | Time for KV caching (CPU-GPU)              |

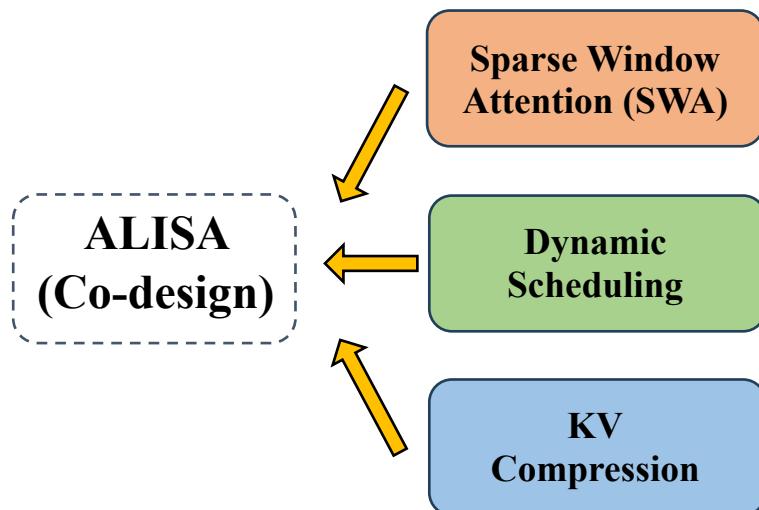
### Solution:

- Divide the problem into two sub-problems, including an I/O problem and a computation problem.
  - For I/O problem, we can use greedy search
  - For computation, we can use profiling

# Algorithm-System Co-Design

## Additional System Optimization

- KV Compression (quantizing FP16 KV tensors to INT8)



**Table I: Comparison of our ALISA and prior works.**

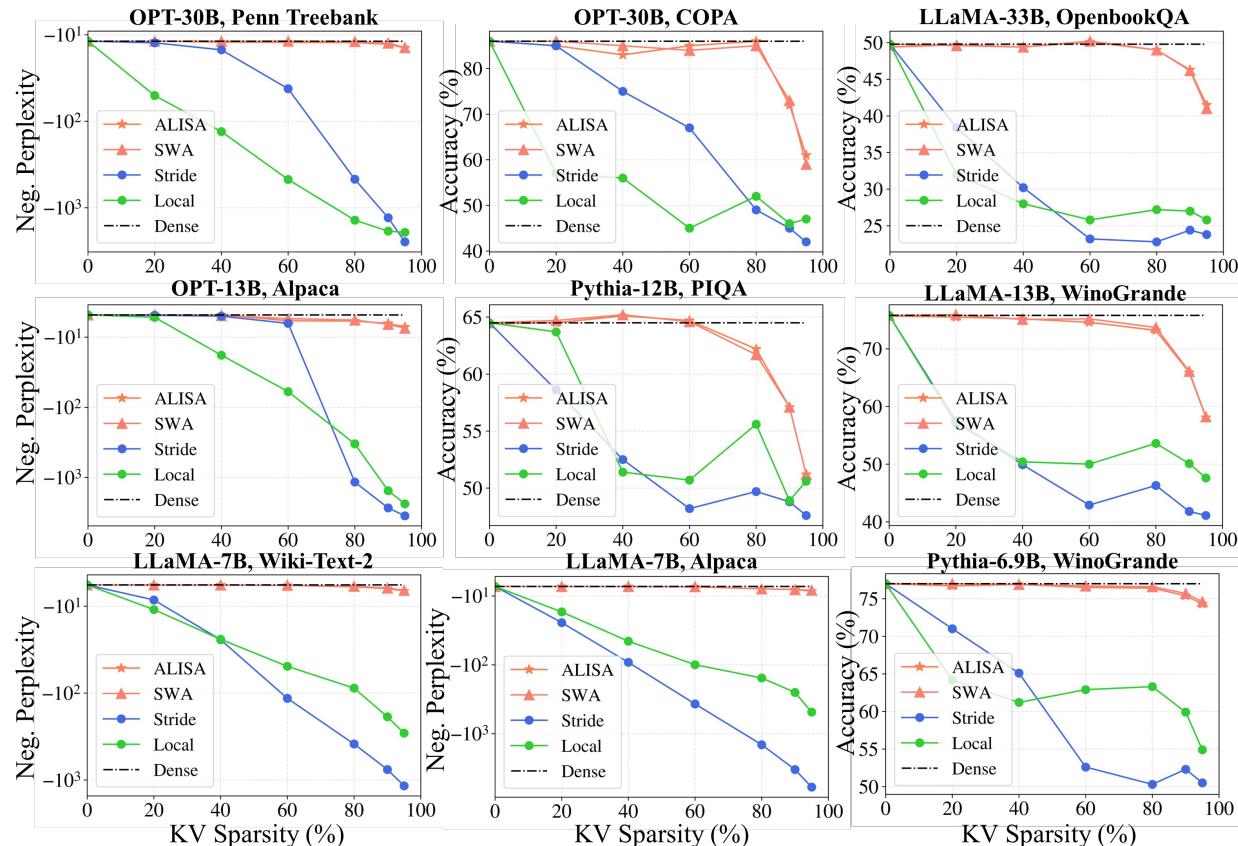
| Design              | vLLM [19]            | FlexGen [29]         | ALISA (Ours)          |
|---------------------|----------------------|----------------------|-----------------------|
| Sparse Attn.        | ✗                    | ✗                    | ✓                     |
| Caching Granularity | Block-level (Static) | Head-level (Static)  | Token-level (Dynamic) |
| Recomputation       | ✓                    | ✗                    | ✓                     |
| Scenario            | Online (Multi-GPU)   | Offline (Single-GPU) | Offline (Single-GPU)  |
| Co-Design           | ✗                    | ✗                    | ✓                     |

# Evaluation

## Experimental Settings:

- Models: OPT (6.7B, 13B, 30B), LLaMA (7B, 13B, 33B), Pythia (6.9B, 12B)
- Algorithm Baselines: Dense attention, Local attention [7], Strided attention [8]
- System Baselines: DeepSpeed [13], Accelerate [14], FlexGen [3], vLLM [4]
- Datasets: Alpaca, Penn Treebank, Wiki-Text-2 (language modeling)  
OpenBookQA, PIQA, COPA, Winogrande (question answering)
- Metrics: Perplexity, Accuracy (Algorithm), Throughput (system)
- Hardware Platforms: V100 16/32 GB, H100 80 GB, 128 GB DRAM (single GPU-CPU sys)

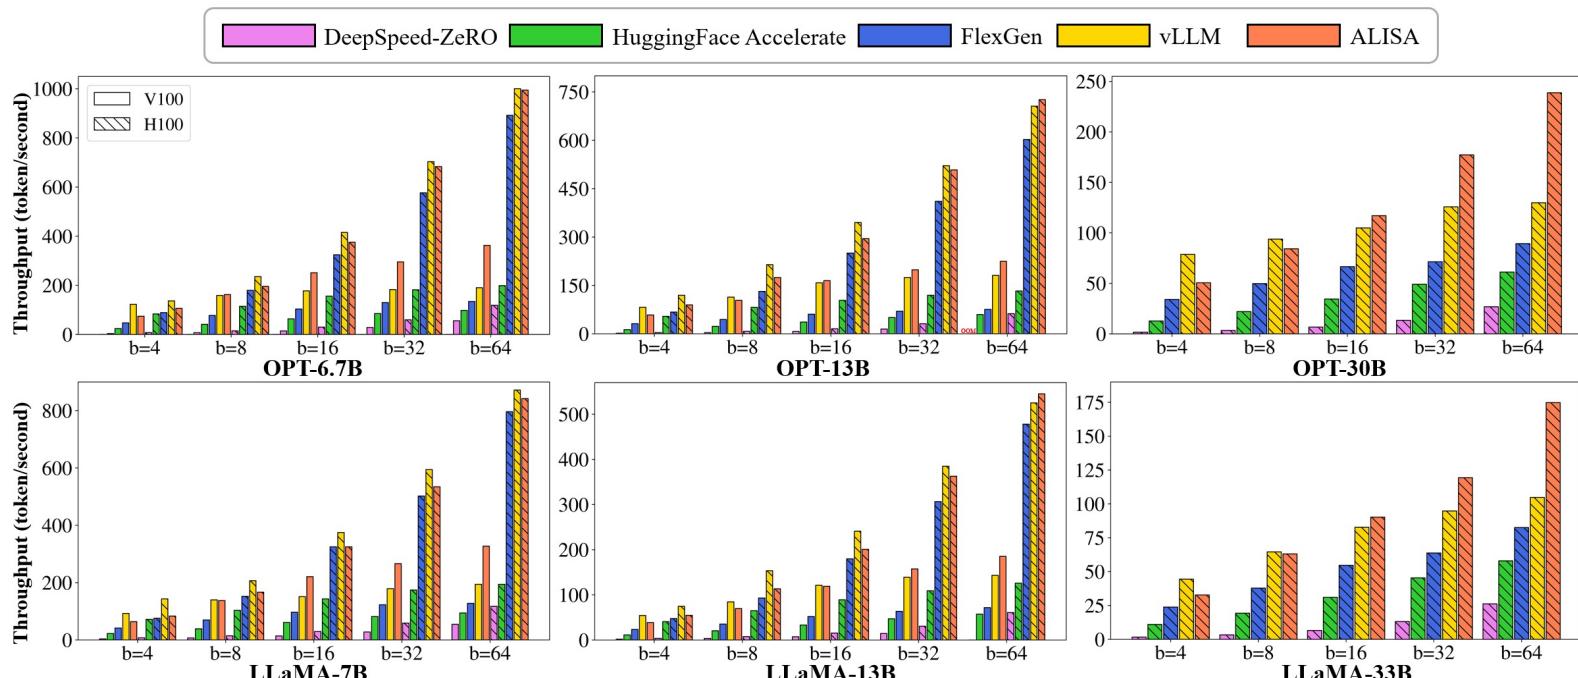
# Algorithm Results



1. ALISA consistently outperforms local and strided attention **across different model types and scales**
2. ALISA can maintain **identical** performance as dense attention to up to **80% KV sparsity**
3. KV compression has almost **no effect** on accuracy

**Full results can be found in the paper**

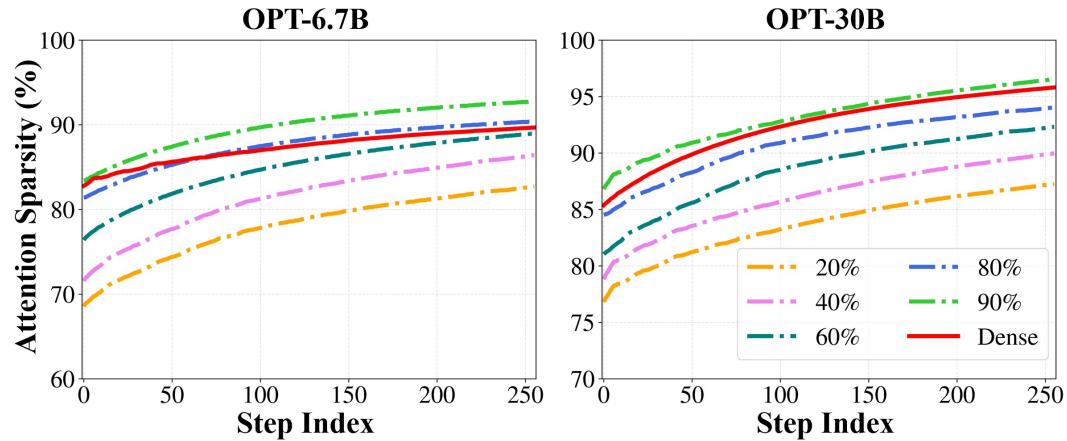
# System Results



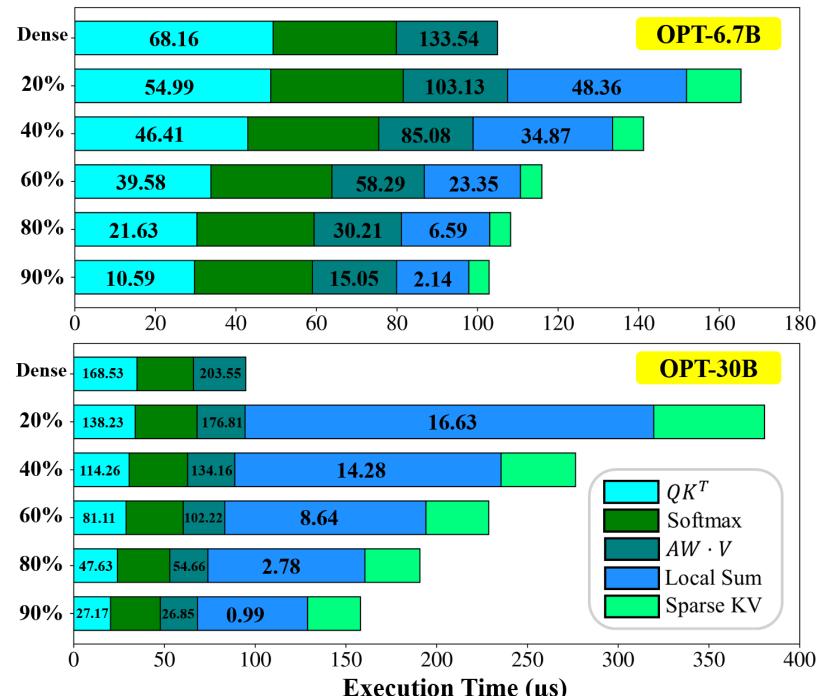
- Compared to FlexGen, ALISA achieves  $1.4\sim3.0\times$  throughput improvement, showing much better scalability across different model sizes and batch sizes
- Compared to vLLM, under large batch sizes, ALISA can sustain up to  $1.9\times$  higher throughput

# Performance Analysis

## Attainable Sparsity



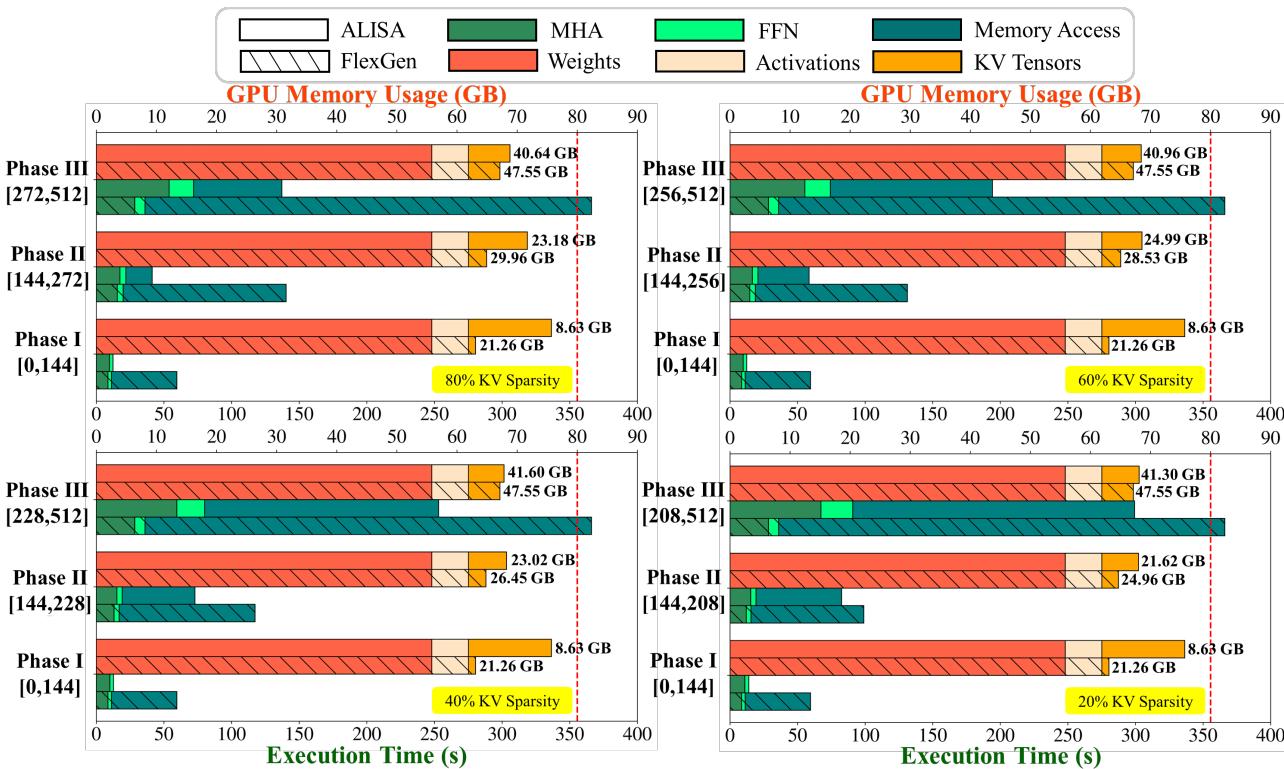
## Kernel-level Breakdown



1. KV Sparsity increases the sparsity by creating sparse KV tensors, which is close to dense attention
2. There exists compute under-utilization in SWA calculation, but the overall overhead is relatively small against dense attention

# Performance Analysis

## LLM Inference Breakdown

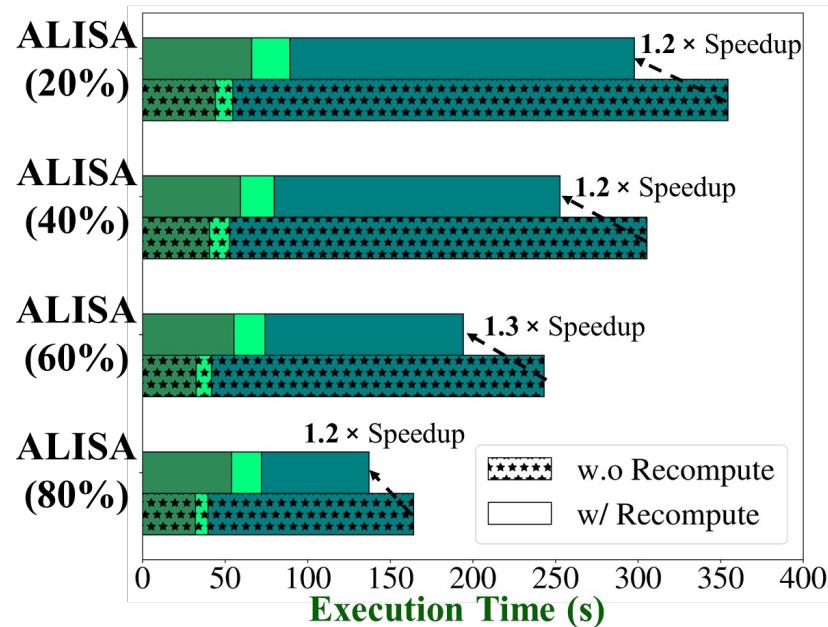


1. ALISA improves upon FlexGen across **different phases**
2. With **higher KV sparsity**, the speedup of ALISA over FlexGen is **more significant**
3. ALISA makes **better use of the GPU memory**

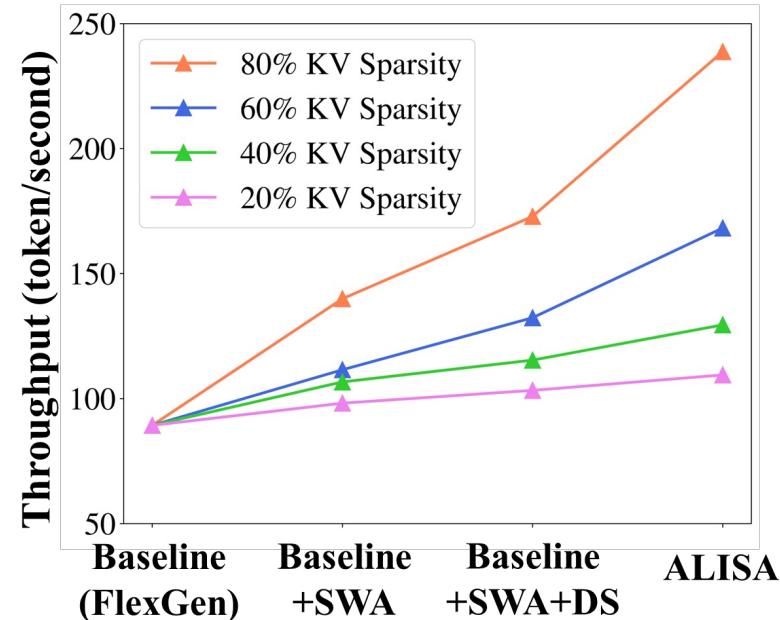
# Performance Analysis

## Ablation Study

### - Impact of Recomputation



### - Impact of each technique



# Conclusion

- We identify the challenges in KV caching for LLM inference and propose **an algorithm-system co-design solution, ALISA**, for efficient LLM inference
- On the algorithm level, we propose **sparse window attention (SWA)** that creates a mixture of globally dynamic and locally static sparse patterns in KV tensors to reduce the memory footprint while maintaining high accuracy.
- On the system level, we design a **three-phase scheduler** to dynamically allocate KV tensors between GPU and CPU memory to reduce data transfer at the token level.
- Extensive experiments demonstrate that ALISA can significantly reduce the memory footprint of KV tensors and increase the throughput over previous baselines, with negligible accuracy drop

# References

- [1] B. Wang. Timeline of Open and Proprietary Large Language Models.
- [2] M. Ott, et al. “Fairseq: A fast, extensible toolkit for sequence modeling,” NACCL, 2019.
- [3] Y. Sheng, et al. “High-throughput generative inference of large language models with a single gpu,” ICML 2023.
- [4] W. Kwon, et al. “Efficient memory management for large language model serving with pagedattention,” SOSP, 2023.
- [5] N. Kitaev , et al. “Reformer: The efficient transformer,” ICLR,.
- [6] S. Wang, et al. “Linformer: Selfattention with linear complexity,” ArXiv, 2020.
- [7] I. Beltagy, et al. “Longformer: The long document transformer,” ArXiv, 2020.
- [8] R. Child, et al. “Generating long sequences with sparse transformers,” ArXiv, 2019.
- [9] H. Wang, et al. “Spatten: Efficient sparse attention architecture with cascade token and head pruning,” HPCA, 2020.
- [10] J. Dass, et al. “Vitality: Unifying low-rank and sparse approximation for vision transformer acceleration with a linear taylor attention,” HPCA, 2022.
- [11] J. Lin , et al. “Awq: Activation-aware weight quantization for llm compression and acceleration,” ArXiv, 2023.
- [12] E. Frantar, et al, “Gptq: Accurate post-training quantization for generative pre-trained transformers,” ICLR, 2023.
- [13] R. Y. Aminabadi, et al. “Deepspeedinference: Enabling efficient inference of transformer models at unprecedented scale,” SC, 2022.
- [14] T. Wolf , et al. “Huggingface’s transformers: State-of-the-art natural language processing,” ArXiv, 2019.

# Acknowledgements

- We appreciate the generous support from the National Science Foundation (NSF) grant 1907765, 2400014
- Faculty and student investigators from the UCF Computer Systems and Data Science (CASS) Laboratory (Youpeng Zhao, Di Wu and many others)

**Thank you!**